

## ← Notes



## ▲ Mo's algorithm

192 Mo's algorithm

Mo algorithm

Algorithm

Code Monk

CodeMonk

## Introduction

Mo's algorithm is a generic idea. It applies to the following class of problems:

You are given array Arr of length N and Q queries. Each query is represented by two numbers L and R, and it asks you to compute some function Func with subarray Arr[L..R] as its argument.

For the sake of brevity we will denote Func([L, R]) as the value of Func on subarray Arr[L..R]. If this sounds too abstract, let's look at specific example:

There is an integer array Arr of length N and Q queries. For each i, query #i asks you to output the sum of numbers on subarray  $[L_i, R_i]$ , i.e.  $Arr[L_i] + Arr[L_i + 1] + ... + Arr[R_i]$ .

Here we have Func([L, R]) = Arr[L] + Arr[L + 1] + ... + Arr[R].

This does not sound so scary, does it? You've probably heard of solutions to this problem using Segment Trees or Binary Indexed Trees, or even prefix sums.

Mo's algorithm provides a way to answer all queries in O((N + Q) \* sqrt(N) \* F) time with at least O(Q) additional memory. Meaning of F is explained below.

The algorithm is applicable if all following conditions are met:

- 1. Arr is not changed by queries;
- 2. All queries are known beforehand (techniques requiring this property are often called "offline algorithms");
- 3. If we know Func([L, R]), then we can compute Func([L + 1, R]), Func([L 1, R]), Func([L, R + 1]) and Func([L, R - 1]), each in **O(F)** time.

Due to constraints on array immutability and queries being known, Mo's algorithm is inapplicable most of the time. Thus, knowing the method can help you on rare occasions. But. Due to property #3, the algorithm can solve problems that are unsolvable otherwise. If the problem was meant to be solved using Mo's algorithm, then you can be 90% sure that it can not be accepted without knowing it. Since the approach is not well-known, in situations where technique is appropriate, you will easily overcome majority of competitors.

## **Basic overview**

We have Q queries to answer. Suppose we answer them in order they are asked in the following manner:

```
for i = 0..Q-1:
   L, R = query #i
   for j = L..R:
    do some work to compute Func([L, R])
```

This can take  $\Omega(N * Q)$  time. If N and Q are of order  $10^5$ , then this would lead to time limit exceeded.

But what if we answer queries in different order? Can we do better then?

#### Definition #1:

**Segment [L, R]** is a continuous subarray Arr[L..R], i.e. array formed by elements Arr[L], Arr[L + 1], ..., Arr[R]. We call L **left endpoint** and R **right endpoint** of segment [L, R]. We say that index i belongs to segment [L, R] if  $L \le i \le R$ .

### Notation

- 1. Throughout this tutorial " $\frac{x}{y}$ " will mean "integer part of x divided by y". For instance,  $\frac{10}{4} = 2$ ,  $\frac{15}{3} = 5$ ,  $\frac{27}{8} = 3$ ;
- 2. By "sqrt(x)" we will mean "largest integer less or equal to square root of x". For example, sqrt(16) = 4, sqrt(39) = 6;
- 3. Suppose a query asks to calculate Func([L, R]). We will denote this query as [L, R] the same way as the respective argument to Func;
- 4. Everything is 0-indexed.

We will describe Mo's algorithm, and then prove its running time.

The approach works as follows:

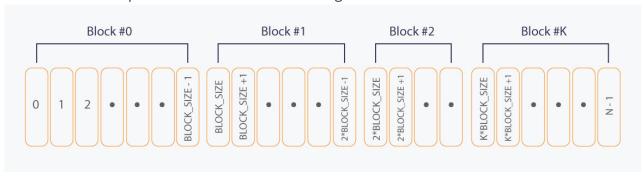
- 1. Denote BLOCK\_SIZE = sqrt(N);
- 2. Rearrange all queries in a way we will call "Mo's order". It is defined like this:  $[L_1, R_1]$  comes earlier than  $[L_2, R_2]$  in Mo's order if and only if:
  - a) L<sub>1</sub>/<sub>BLOCK\_SIZE</sub> < L<sub>2</sub>/<sub>BLOCK\_SIZE</sub>
  - b)  $L_{BLOCK\ SIZE} == L_{BLOCK\ SIZE} \&\& R_1 < R_2$
- 3. Maintain segment [mo\_left, mo\_right] for which we know Func([mo\_left, mo\_right]). Initially, this segment is empty. We set mo\_left = 0 and mo\_right = -1;

- 4. Answer all queries following Mo's order. Suppose the next query you want to answer is [L, R]. Then you perform these steps:
  - a) while mo\_right is less than R, extend current segment to [mo\_left, mo\_right + 1];
  - b) while mo\_right is greater than R, cut current segment to [mo\_left, mo\_right 1];
  - c) while mo\_left is greater than L, extend current segment to [mo\_left 1, mo\_right];
  - d) while mo\_left is less than L, cut current segment to [mo\_left + 1, mo\_right].

This will take O( (|| eft - L| + | right - R|) \* F) time, because we required that each extension\deletion is performed in O(F) steps. After all transitions, you will have mo\_left = L and mo\_right = R, which means that you have successfully computed Func([L, R]).

## Time complexity

Let's view Arr as a union of disjoint segments of size BLOCK\_SIZE, which we will call "blocks". Take a look at the picture for better understanding:



Let K be the index of last block. Then there are K + 1 blocks, because we number them from zero. Notice than K'th block can have less than BLOCK\_SIZE elements.

Proposition #1:

K = O(sqrt(N)).

Proof:

If sqrt(N) \* sqrt(N) = N (which may be false due to our definition of sqrt(N)), then K = sqrt(N) - 1, because we have K + 1 blocks, each of size sqrt(N). Otherwise, K = sqrt(N), because we need one additional block to store N - sqrt(N) \* sqrt(N) elements.

**UPD 19 june 2017**: As pointed out in the comments, this does not hold for small values of N. Alternative proof (less intuitive, more formal): we need the smallest K such that  $K * \operatorname{sqrt}(N) \ge N$  - when this is true, our space is sufficient to hold all N elements. If we divide by  $\operatorname{sqrt}(N)$ , we get  $K \ge N / \operatorname{sqrt}(N)$ .  $\operatorname{sqrt}(N)$  is either  $\operatorname{real\_sqrt}(N)$  (the one without truncating value after decimal point) or  $\operatorname{real\_sqrt}(N)$  - 1. So  $K \ge N / \operatorname{(real\_sqrt}(N) - 1)$ .

constant (you can find it, for example, by taking the first 1000 values of N and taking maximum value of the bound). So K - real\_sqrt(N)  $\leq$  constant, which implies K = O(sqrt(N)).

## Proposition #2:

Block #r is a segment [r \* BLOCK\_SIZE, min(N - 1, (r + 1) \* BLOCK\_SIZE - 1)].

Proof (by induction):

For block #0 statement is true — it is a segment [0, BLOCK\_SIZE - 1], containing BLOCK\_SIZE elements.

Suppose first  $T \le K$  blocks satisfy the above statement. Then the last of those blocks is a segment  $[(T - 1) * BLOCK\_SIZE, T * BLOCK\_SIZE - 1]$ .

Then we form the next, T+1'th block. First element of this block will have array index T \*  $BLOCK\_SIZE$ . We have at most  $BLOCK\_SIZE$  elements to add to block (there may be less if T + 1 = K + 1). So the last index in T+1'th block will be min(N - 1, (T + 1) \*  $BLOCK\_SIZE$  - 1).

Corollary #1: Two indices i and j belong to same block #r if and only if  $i_{BLOCK\_SIZE} = i_{BLOCK\_SIZE} = r$ .

### Definition #2:

 $Q_r = \{ \text{ query } [L, R] \mid L'_{BLOCK\_SIZE} = r \}$ . Informally,  $Q_r$  is a set of queries from input, whose left endpoints belong to block #r. Notice that this set may be empty. We denote the size of  $Q_r$  as  $|Q_r|$ .

## Proposition #3:

For each r, queries from  $Q_r$  lie continuously in Mo's order and they appear in it in non-decreasing order of right endpoints.

Queries from  $Q_r$  come earlier than queries from  $Q_{r+1}$  for every r = 0..K-1.

*Proof* follows from definition of Mo's order.

Corollary #2: When we are processing queries following Mo's order, we firstly process all queries from  $Q_0$ , then all queries from  $Q_1$ , and so on, up to  $Q_K$ .

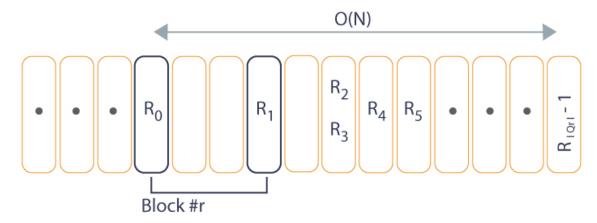
## Theorem #1:

Mo\_right changes its value O(N \* sqrt(N)) times throughout the run of Mo's algorithm. *Proof:* 

1. Suppose we've just started processing queries from  $Q_r$  for some r, and we've already answered first (following Mo's order) query from it. Let that query be [L,  $R_0$ ]. This means that now mo\_right =  $R_0$ .

Let  $R_0$ ,  $R_1$ , ...,  $R_{|Q_r|-1}$  be right endpoints of queries from  $Q_r$ , in order they appear in Mo's order. From proposition #3 we know that

$$R_0 \le R_1 \le R_2 \le ... \le R_{|Q_r|-1}$$



From proposition #2 and definition of Q<sub>r</sub> we know that

r \* BLOCK SIZE ≤ L

Since right endpoint is not lower that left endpoint (i.e.  $L \le R$ ), we conclude that  $r * BLOCK\_SIZE \le R_0$ 

We have N elements in Arr, so any query has right endpoint less than N. Therefore,

$$R_{|O_r|-1} \le N-1$$

Since R's are not decreasing, the total amount of times mo\_right changes is  $R_{|O_r|-1} - R_0 \le N - 1 - r * BLOCK_SIZE = O(N)$ 

(we've substituted  $R_{|Q_r|-1}$  with its highest possible value, and  $R_0$  with its lowest possible value to maximize the subtraction).

There are O(sqrt(N)) different values of r (proposition #1), so in total we will have O(N \* sqrt(N)) changes, assuming that we've already started from the first query of each  $Q_r$ .

2. Now suppose we've just ended processing queries from  $Q_r$  and we must process first query from  $Q_{r+1}$  (assuming that it is not empty. If it is, then choose next non-empty set). Currently, mo\_right is constrained to be:

r \* BLOCK\_SIZE  $\leq$  mo\_right  $\leq$  N - 1

Let the first query from  $Q_{r+1}$  be [L', R']. We know (similarly to previous paragraph) that  $(r+1)*BLOCK\_SIZE \le R' \le N-1$ 

Hence, mo\_right must be changed at most

times (we took lowest value of mo\_right with highest value of R' and vice-versa). This is clearly O(N) (and it does not matter whether it is r+1'th set of r+k'th for some k > 1). There are O(sqrt(N)) switches from r to r + 1 (and it's true even if we skip some empty sets), so in total we will have O(N \* sqrt(N)) mo\_right changes to do this.

There are no more cases when mo\_right changes. Overall, we have O(N \* sqrt(N)) + O(N \* sqrt(N)) = O(N \* sqrt(N)) changes of mo\_right.

Corollary #3: All mo\_right changes combined take O(N \* sqrt(N) \* F) time (because each change is done in O(F) time).

### Theorem #2:

Mo\_left changes its value O(Q \* sqrt(N)) times throughout the run of Mo's algorithm. *Proof:* 

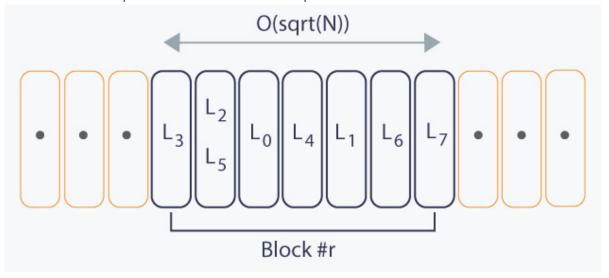
1. Suppose, as in the proof of Theorem #1, we've just started processing queries from  $Q_r$  for some r, and currently mo\_left = L, mo\_right =  $R_0$ , where [L,  $R_0$ ]  $\in Q_r$ . For every query [L', R']  $\in Q_r$  by definition #2 and proposition #2 we have:

$$r * BLOCK_SIZE \le L' \le (r + 1) * BLOCK_SIZE - 1$$

So, when we change mo\_left from one query to other (both queries  $\in Q_r$ ), we must do at most

$$(r + 1) * BLOCK_SIZE - 1 - r * BLOCK_SIZE = BLOCK_SIZE - 1 = O(sqrt(N))$$
 changes to mo\_left. At the picture below we represented leftpoints that lie in block #r,

where the subscript shows relative order of queries:



There are  $|Q_r|$  queries in  $Q_r$ . That means, that we can estimate upper bound on number of changes for single r as  $O(|Q_r| * \text{sqrt}(N))$ . Let's sum it over all r:

$$O(sqrt(N) * (|Q_0| + |Q_1| + ... + |Q_K|)) = O(sqrt(N) * Q)$$

(because each query's leftpoint belongs to exactly one block, and we have Q queries in total).

2. Now suppose we're done with queries from  $Q_r$  and want to process queries from nonempty set  $Q_{r+k}$  (for some k > 0). Any query  $[L', R'] \in Q_{r+k}$  has

$$(r + k) * BLOCK_SIZE \le L' \le (r + k + 1) * BLOCK_SIZE - 1$$

Similarly, any query [L, R]  $\in$  Q<sub>r</sub> has

$$r * BLOCK SIZE \le L \le (r + 1) * BLOCK SIZE - 1$$

We can see that the maximum number of changes needed to transit from some query from  $Q_r$  to some query from  $Q_{r+k}$  is

Now if we sum this over all r, we will get at most

```
K * BLOCK\_SIZE = O(sqrt(N)) * O(sqrt(N)) = O(N), because sum of all possible k's does not exceed K.
```

There are no more cases when mo\_left changes. Overall, we have  $O(\operatorname{sqrt}(N) * Q) + O(N) = O(\operatorname{sqrt}(N) * Q)$  changes of mo\_left.

Corollary #4: All mo\_left changes combined take O(Q \* sqrt(N) \* F) time (because each change is done in O(F) time).

Corollary #5: Time complexity of Mo's algorithm is O((N + Q) \* sqrt(N) \* F).

## **Example problem**

Let's look at an example problem (idea taken from here):

You have an array Arr of N numbers ranging from 0 to 99. Also you have Q queries [L, R]. For each such query you must print

$$V([L, R]) = \sum_{i=0...99} count(i)^2 * i$$

where count(i) is the number of times i occurs in Arr[L..R].

Constraints are  $N \le 10^5$ ,  $Q \le 10^5$ .

To apply Mo's algorithm, you must ensure of three properties:

- 1. Arr is not modified by queries;
- 2. Queries are known beforehand;
- 3. If you know V([L, R]), then you can compute V([L + 1, R]), V([L 1, R]), V([L, R 1]) and V([L, R + 1]), each in O(F) time.

First two properties obviously hold. The third property depends on the time bound — O(F). Surely, we can compute V([L + 1, R]) from scratch in  $\Omega(R - L) = \Omega(N)$  in the worst case. But looking at complexity of Mo's algorithm, you can deduce that this will surely time out, because we multiply O(F) with O((Q + N) \* sqrt(N)).

Typically, you want O(F) to be O(1) or O(log(n)). Choosing a way to achieve appropriate time bound O(F) is programmer's main concern when solving problems using Mo's algorithm.

I will describe a way to achieve O(F) = O(1) for this problem.

Let's maintain variable **current\_answer** (initialized by zero) to store V([mo\_left, mo\_right]) and integer array **cnt** of size 100, where cnt[x] will be the number of times x occurs in [mo\_left, mo\_right].

From the definition of current\_answer we can see that

```
current answer = V([mo\_left, mo\_right]) = \sum_{i=0.99} count'(i)^2 * i
```

where count'(x) is the number of times x occurs in [mo\_left, mo\_right].

```
By the definition of cnt we see that count'(x) = cnt[x], so current answer = \sum_{i=0..99} cnt[i]<sup>2</sup> * i.
```

Suppose we want to change mo\_right to mo\_right + 1. It means we want to add number p = Arr[mo\_right + 1] into consideration. But we also want to retain current\_answer and cnt's properties.

Maintaining cnt is easy: just increase cnt[p] by 1. It is a bit trickier to deal with current\_answer.

We know (from the definitions) that current\_answer contains summand  $cnt[p]^2 * p$  before addition. Let's subtract this value from the answer. Then, after we perform addition to cnt, add again  $cnt[p]^2 * p$  to the answer (make no mistake: this time it will contain updated value of cnt[p]). Both updates take O(1) time.

All other transitions (mo\_left to mo\_left + 1, mo\_left to mo\_left - 1 and mo\_right to mo\_right - 1) can be done in the same way, so we have O(F) = O(1). You can refer to the code below for clarity.

Now let's look into detail on one sample test case:

## Input:

```
Arr = [0, 1, 1, 0, 2, 3, 4, 1, 3, 5, 1, 5, 3, 5, 4, 0, 2, 2] of 18 elements

Queries (0-indexed): [0, 8], [2, 5], [2, 11], [16, 17], [13, 14], [1, 17], [17, 17]
```

The algorithm works as follows:

Firstly, set BLOCK\_SIZE = sqrt(18) = 4. Notice that we have 5 blocks: [0, 3], [4, 7], [8, 11], [12, 15], [16, 17]. The last block contains less than BLOCK SIZE elements.

Then, set  $mo_left = 0$ ,  $mo_right = -1$ ,  $current_answer = 0$ , cnt = [0, 0, 0, 0, 0, 0] (I will use only first 6 elements out of 100 for the sake of simplicity).

Then sort queries. The Mo's order will be:

[2,5], [0, 8], [2, 11], [1, 17] (here ends  $Q_0$ ) [13, 14] (here ends  $Q_3$ ) [16, 17], [17, 17] (here ends  $Q_4$ ).

Now, when everything is set up, we can answer queries:

1. We need to process query [2, 5]. Currently, our segment is [0, -1]. So we need to move mo right to 5 and mo left to 2.

```
Let's move mo_right first:
```

```
mo_right = 0, current_answer = 0, cnt = [1, 0, 0, 0, 0, 0]
mo_right = 1, current_answer = 1, cnt = [1, 1, 0, 0, 0, 0]
mo_right = 2, current_answer = 4, cnt = [1, 2, 0, 0, 0, 0]
mo_right = 3, current_answer = 4, cnt = [2, 2, 0, 0, 0, 0]
mo_right = 4, current_answer = 6, cnt = [2, 2, 1, 0, 0, 0]
mo_right = 5, current_answer = 9, cnt = [2, 2, 1, 1, 0, 0]
Now we must move mo_left:
mo_left = 1, current_answer = 9, cnt = [1, 2, 1, 1, 0, 0]
```

```
mo_left = 2, current_answer = 6, cnt = [1, 1, 1, 1, 0, 0]
Thus, the answer for query [2, 5] is 6.
```

2. Our next query is [0, 8]. Current segment [mo\_left, mo\_right] is [2, 5]. We need to move mo\_right to 8 and mo\_left to 0.

```
Again, let's move mo_right first:
```

```
mo_right = 6, current_answer = 10, cnt = [1, 1, 1, 1, 1, 0]
```

Then we move mo\_left:

So, the answer for query [0, 8] is 27.

3. Next query is [2, 11]. Current segment is [0, 8]. We need to move mo\_right to 11 and mo\_left to 2.

```
mo_right = 9, current_answer = 32, cnt = [2, 3, 1, 2, 1, 1]
```

Answer for query [2, 11] is 47.

4. Next query is [1, 17]. Current segment is [2, 11]. We need to move mo\_right to 17 and mo\_left to 1.

```
mo_right = 12, current_answer = 62, cnt = [1, 3, 1, 3, 1, 2]
```

Answer for query [1, 17] is 122.

5. Our next goal is query [13, 14]. Notice that it starts in different block from the previous query [1, 17]. Consequently, this is the first time mo\_right will move to the left. We need to move mo\_left to 13 and mo\_right to 14.

```
mo_right = 16, current_answer = 112, cnt = [2, 4, 2, 3, 2, 3]
```

```
mo_left = 8, current_answer = 62, cnt = [0, 1, 0, 2, 1, 3]
mo_left = 9, current_answer = 53, cnt = [0, 1, 0, 1, 1, 3]
mo_left = 10, current_answer = 28, cnt = [0, 1, 0, 1, 1, 2]
mo_left = 11, current_answer = 27, cnt = [0, 0, 0, 1, 1, 2]
mo_left = 12, current_answer = 12, cnt = [0, 0, 0, 1, 1, 1]
mo_left = 13, current_answer = 9, cnt = [0, 0, 0, 0, 1, 1]
Answer for query [13,14] is 9.
```

6. Next query is [16, 17]. Notice, however, that now we do not need to move mo\_right to the left. We need to move mo\_left to 16 and mo\_right to 17.

```
mo_right = 15, current_answer = 9, cnt = [1, 0, 0, 0, 1, 1]
mo_right = 16, current_answer = 11, cnt = [1, 0, 1, 0, 1, 1]
mo_right = 17, current_answer = 17, cnt = [1, 0, 2, 0, 1, 1]
mo_left = 14, current_answer = 12, cnt = [1, 0, 2, 0, 1, 0]
mo_left = 15, current_answer = 8, cnt = [1, 0, 2, 0, 0, 0]
mo_left = 16, current_answer = 8, cnt = [0, 0, 2, 0, 0, 0]
Answer for [16, 17] is 8.
```

7. The last query is [17, 17]. It requires us to move mo\_left one unit to the right.

mo\_left = 17, current\_answer = 2, cnt = [0, 0, 1, 0, 0, 0]

Answer for this query is 2.

Now the important part comes: we must **output answers** not in order we've achieved them, but in **order they were asked**.

### **Output:**

27

6

47

8

9

122

2

## **Implementation**

Here is the C++ implementation for the above problem:

```
#include <bits/stdc++.h>
using namespace std;

int N, Q;

// Variables, that hold current "state" of computation
long long current_answer;
```

```
long long cnt[100];
// Array to store answers (because the order we achieve them is messed up)
long long answers[100500];
int BLOCK SIZE;
int arr[100500];
// We will represent each query as three numbers: L, R, idx. Idx is
// the position (in original order) of this query.
pair< pair<int, int>, int> queries[100500];
// Essential part of Mo's algorithm: comparator, which we will
// use with std::sort. It is a function, which must return True
// if query x must come earlier than query y, and False otherwise.
inline bool mo_cmp(const pair< pair<int, int>, int> &x,
         const pair< pair<int, int>, int> &y)
{
    int block x = x.first.first / BLOCK SIZE;
    int block y = y.first.first / BLOCK SIZE;
    if(block x != block y)
         return block_x < block_y;</pre>
    return x.first.second < y.first.second;</pre>
}
// When adding a number, we first nullify it's effect on current
// answer, then update cnt array, then account for it's effect again.
inline void add(int x)
{
    current_answer -= cnt[x] * cnt[x] * x;
    cnt[x]++;
    current_answer += cnt[x] * cnt[x] * x;
}
// Removing is much like adding.
inline void remove(int x)
{
    current_answer -= cnt[x] * cnt[x] * x;
    cnt[x]--;
    current answer += cnt[x] * cnt[x] * x;
}
```

```
int main()
{
    cin.sync with stdio(false);
    cin >> N >> Q;
    BLOCK SIZE = static cast<int>(sqrt(N));
    // Read input array
    for(int i = 0; i < N; i++)</pre>
         cin >> arr[i];
    // Read input queries, which are 0-indexed. Store each query's
    // original position. We will use it when printing answer.
    for(int i = 0; i < Q; i++) {</pre>
         cin >> queries[i].first.first >> queries[i].first.second;
         queries[i].second = i;
    }
    // Sort queries using Mo's special comparator we defined.
    sort(queries, queries + Q, mo cmp);
    // Set up current segment [mo_left, mo_right].
    int mo_left = 0, mo_right = -1;
    for(int i = 0; i < Q; i++) {</pre>
         // [left, right] is what query we must answer now.
         int left = queries[i].first.first;
         int right = queries[i].first.second;
         // Usual part of applying Mo's algorithm: moving mo_left
         // and mo_right.
         while(mo_right < right) {</pre>
             mo_right++;
              add(arr[mo_right]);
         while(mo_right > right) {
              remove(arr[mo_right]);
             mo_right--;
         }
         while(mo left < left) {</pre>
              remove(arr[mo left]);
             mo_left++;
```

```
while(mo_left > left) {
          mo_left--;
          add(arr[mo_left]);
}

// Store the answer into required position.
     answers[queries[i].second] = current_answer;
}

// We output answers *after* we process all queries.
for(int i = 0; i < 0; i++)
          cout << answers[i] << "\n";
return 0;
}</pre>
```

Same solution without global variables (the way I like to implement it):

```
#include <bits/stdc++.h>
using std::vector;
using std::tuple;
 * Take out adding\removing logic into separate class.
 * It provides functions to add and remove numbers from
 * our structure, while maintaining cnt[] and current_answer.
 */
struct Mo
{
    static constexpr int MAX_VALUE = 1005000;
    vector<long long> cnt;
    long long current_answer;
    void process(int number, int delta)
    {
         current answer -= cnt[number] * cnt[number] * number;
         cnt[number] += delta;
         current answer += cnt[number] * cnt[number] * number;
    }
public:
    Mo()
```

```
{
         cnt = vector<long long>(MAX_VALUE, 0);
         current answer = 0;
    }
    long long get_answer() const
         return current_answer;
    }
    void add(int number)
         process(number, 1);
    }
    void remove(int number)
         process(number, -1);
};
int main()
{
    int N, Q, BLOCK_SIZE;
    std::cin.sync_with_stdio(false);
    std::cin >> N >> Q;
    BLOCK_SIZE = static_cast<int>(sqrt(N));
    // No global variables, everything inside.
    vector<int> arr(N);
    vector<long long> answers(Q);
    vector< tuple<int, int, int> > queries;
    queries.reserve(Q);
    for(int i = 0; i < N; i++)</pre>
         std::cin >> arr[i];
    for(int i = 0; i < Q; i++) {</pre>
         int le, rg;
         std::cin >> le >> rg;
         queries.emplace_back(le, rg, i);
    }
```

```
// Comparator as a lambda-function, which catches BLOCK_SIZE
// from the above definition.
auto mo cmp = [BLOCK SIZE](const tuple<int, int, int> &x,
         const tuple<int, int, int> &y) -> bool {
    int block x = std::get<0>(x) / BLOCK SIZE;
    int block_y = std::get<0>(y) / BLOCK_SIZE;
    if(block_x != block_y)
         return block_x < block_y;</pre>
    return std::get<1>(x) < std::get<1>(y);
};
std::sort(queries.begin(), queries.end(), mo_cmp);
Mo solver;
int mo_left = 0, mo_right = -1;
// Usual Mo's algorithm stuff.
for(const auto &q: queries) {
    int left, right, answer_idx;
    std::tie(left, right, answer_idx) = q;
    while(mo_right < right) {</pre>
         mo_right++;
         solver.add(arr[mo right]);
    while(mo_right > right) {
         solver.remove(arr[mo right]);
         mo_right--;
    }
    while(mo_left < left) {</pre>
         solver.remove(arr[mo left]);
         mo left++;
    while(mo_left > left) {
         mo_left--;
         solver.add(arr[mo_left]);
    }
    answers[answer idx] = solver.get answer();
}
```

```
for(int i = 0; i < Q; i++)</pre>
          std::cout << answers[i] << "\n";</pre>
     return 0;
}
```

# **Practice problems**

In case you want to try out implementing Mo's technique for yourself, check out these problems:

## 1. Kriti and her birthday gift

Difficulty: easy.

Prerequisites: string hashing.

Comment: tests for this problem are flawed (you can get at most 33 out of 100 for this problem), because both setter's and tester's implementations from the editorial have the same bug in hashing function. Can you see what it is?

Reference implementation: here.

### 2. SUBSTRINGS COUNT

Difficulty: easy.

Prerequisites: string hashing.

Comment: almost the same as previous problem, but this one asks slightly different question and has well-formed tests.

Reference implementation: here.

## 3. Sherlock and inversions

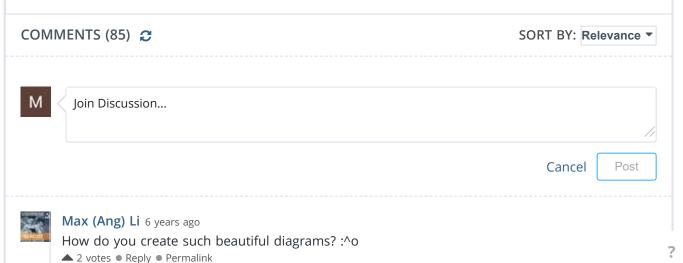
Difficulty: medium.

Prerequisites: segment tree\binary indexed tree, coordinate compression.

Comment: nice and clean problem on Mo's algorithm. Might seem difficult for people not

familiar with prerequisites.

Reference implementation: here.





Mike Koltsov 4 Author 6 years ago

HackerEarth staff did this for me

▲ 8 votes • Reply • Permalink



### Satyarth Agrahari 5 years ago

can you please help me with a sequence to learn algorithms. That will be really helpful. Thanks in advance

▲ 1 vote • Reply • Permalink



Mike Koltsov 4 Author 5 years ago

I think CodeMonk provides reasonable sequence if you don't know where to start. Visit https://www.hackerearth.com/codemonk/

▲ 3 votes • Reply • Permalink



### Satyarth Agrahari 5 years ago

thank you so much

▲ 0 votes • Reply • Permalink



### Sharvin Jondhale 4 years ago

Maybe we should take K=ceil(sqrt(n)), for example if N=45, then sqrt(n)=6, K=6+1=7, available space in array (K+1)\*sqrt(n) = 42, but the array needs to hold 45, so we should take the ceil value instead of floor?

▲ 4 votes • Reply • Permalink



#### Mike Koltsov 4 Author 4 years ago

It does not really matter. You can say "hey, I want at most sqrt(n) in each block". And the number of blocks - that is, our K - will be approximately sqrt(n). You don't need to tune it manually.

▲ 1 vote • Reply • Permalink



#### Sharvin Jondhale 4 years ago

Yeah, you are right, initially I had thought we will be actually splitting them into some "container" ...

▲ 0 votes • Reply • Permalink



### Ankit Kumar 4 years ago

Hi Mike can the last block have more than block size elements. As in the case when n = 8. block size is 2. Therefore K = 2

and block 0 contains 2 elements , block 1 contains 2 elements, block 2 contains 4 elements. Is this true ??

▲ 2 votes • Reply • Permalink



#### Mike Koltsov 4 Author 4 years ago

Hi! Seems like I was wrong in my claim that K is at most sqrt(N). It can be slightly higher, but still O(sqrt(N)).

Last block cannot have more than BLOCK\_SIZE elements: you make another block for that.

▲ 1 vote • Reply • Permalink



#### Pratyush Gaurav Jagaty 4 years ago

Is this "mo\_right to 14" a typo?

Next query is [16, 17]. Notice, however, that now we do not need to move mo\_right to the left. We need to move mo\_left to 16 and mo\_right to 14. mo\_right becomes 17, right?

▲ 1 vote • Reply • Permalink



Mike Koltsov 4 Author 4 years ago

Thank you, I updated my post!

?

▲ 2 votes • Reply • Permalink



## Pratyush Gaurav Jagaty 4 years ago

Thanks, Mike! it was definitely a good read.

▲ 0 votes • Reply • Permalink



#### Aman Goel 6 years ago

Amazing!

I got to learn Fenwick tree as well because of the last problem Thanks a lot for such a nice explanation of the algorithm

▲ 2 votes • Reply • Permalink



#### Arpan Mukherjee 5 years ago

WOW! Really nice post, beautifully explained. Thank you. Just learnt the MO's algorithm. Just have one question.

is Square root decomposition and MO' Algorithm same or different? I mean what's the difference between these two(If they're different).

▲ 0 votes • Reply • Permalink



### Mike Koltsov 4 Author 5 years ago

Thank you for your kind words!

In my opinion, Mo's algorithm is one example of a very broad topic, which is Square root decomposition.

▲ 2 votes • Reply • Permalink



#### Arpan Mukherjee 5 years ago

Okay.. Thank you:)

▲ 0 votes • Reply • Permalink



#### Maddela Sai Karthik 5 years ago

i think there's a typo in the theorem 2 proof part 1 this line:

There are |Qr| queries in Qr. That means, that we can estimate upper bound on number of changes for single r as O(|Qr| \* N). Let's sum it over all r:

O(N \* (|Q0| + |Q1| + ... + |QK|)) = O(sqrt(N) \* Q)

the N should be actually sqrt(N) in the left hand side of the equation..

▲ 1 vote • Reply • Permalink



#### Mike Koltsov 4 Author 5 years ago

Thank you, nice catch!

▲ 0 votes • Reply • Permalink



#### Maddela Sai Karthik 5 years ago

Can you write a similar article for heavy light decomposition the proof Part about it in your free time.

▲ 1 vote • Reply • Permalink



#### Amulya Gaur 3 years ago

Those who are getting TLE in Powerful Arrays problem: Use min number of long long's.. this gave AC in 4.056s

▲ 2 votes • Reply • Permalink



XYZ a year ago

Good one!!

▲ 2 votes • Reply • Permalink

Vijay pandit a year ago



why mo\_left =0 and mo\_right=-1?

▲ 1 vote • Reply • Permalink



Mike Koltsov 4 Author a year ago

Because my thought is as follows: I compute an answer to the segment [mo\_left, mo\_right], with both ends included.

At the beginning, I haven't computed any answer. Hence, I set mo\_right to be less than mo\_left - it denotes an "empty segment".

Instead of that, I could have computed answer for segment [0, 0] - that is, the first element. However, this would lead to more code.

▲ 1 vote • Reply • Permalink



Vijay pandit a year ago

Yep, Thanks for quick reply

▲ 0 votes • Reply • Permalink



#### Arjit Srivastava 6 years ago

Thank you for this amazing post, Mike! :)

▲ 1 vote • Reply • Permalink



#### Viet Nguyenkhanh 6 years ago

thanks you very much, the notes very interesting

▲ 1 vote • Reply • Permalink



#### Anand Hariharan 6 years ago

nice post Mike!

▲ 1 vote • Reply • Permalink



## lago Shavidze 5 years ago

great post,but unfortunately this code gives me time limits on codeforces http://codeforces.com/contest/86/submission/15758953 this is my code,can anyone help me?

▲ 0 votes • Reply • Permalink



#### Mike Koltsov 4 Author 5 years ago

```
Rewriting updates to
inline void upd(ll x, ll mult)
cur_ans += 2 * mult * cnt[x] * x + x;
cnt[x] += mult;
void add(ll x){
upd(x, 1);
}
void removeX(II x){
upd(x, -1);
}
Leads to AC in 4.9sec.
Also you can optimize by memsetting cnt to 0 when you encounter new block (instead of
moving R to the left), like this (taken from my code):
if(i == 0 \text{ or } q[i].L / len != q[i - 1].L / len) { // new bucket}
memset(cnt, 0, sizeof(cnt));
cur = 0;
for(int j = q[i].L; j \le q[i].R; j++) {
// cout << a[i] << " " << cur;
upd(a[j], 1, cur);
// cout << "->" << cur << endl;
```

?

```
L = q[i].L, R = q[i].R;
ans[q[i].i] = cur;
} else {
...
}
• 0 votes • Reply • Permalink
```



#### Sai Teja Utpala 4 years ago

First, most elegant explanation @Mike Koltsov

and To see why this is an optimization i did mathematical analysis(or at least i believe so) following your prints

It goes this way:

if we optimize this way case '2' in both theorems goes away and we get new factor upon which running time depends and that is

worst case update operations we have by making transition to new block and computing first query in it would be

|(r\*block\_size)-(N-1)|, and summing this over all this 'r'

 $=\Sigma | (r*block_size) - (N-1) |$  and 'r' goes from 0 to (k=(sqrt(N)))

= | (sqrt(N)\*block\_size)-(sqrt(N)\*N) |

=O(N\*sqrt(N)-N)

total running time would be O(N\*sqrt(N) + Q\*sqrt(N) + N\*sqrt(N)-N) which better than O(N\*sqrt(N) + N\*sqrt(N + Q\*sqrt(N) + N) by factor of 2N which is crucial for larger N in the problem

correct me if am wrong.?

▲ 0 votes • Reply • Permalink



Mike Koltsov / Author 4 years ago

Well, if you already have N \* sqrt(N) running time, then 2 \* N operations is a small thing.

▲ 0 votes • Reply • Permalink



Mike Koltsov 4 Author 5 years ago

You can refer to my submissions for optimizations mentioned above: http://codeforces.com/contest/86/submission/15762653 (your solution with first optimization, runs in 4.9s)

http://codeforces.com/contest/86/submission/13404487 (my solution, both optimizations, runs in 3.2s)

▲ 0 votes • Reply • Permalink



#### lago Shavidze 5 years ago

thanks a lot, i have one question about this algorithm, i did not quite understand why we must sort in this logic the queries.

▲ 0 votes • Reply • Permalink



Mike Koltsov 4 Author 5 years ago

Because it leads to better runtime complexity, compared to naive approach.

▲ 0 votes • Reply • Permalink



lago Shavidze 5 years ago

thanks, now understand.

▲ 0 votes • Reply • Permalink



shavidze shavidze 5 years ago

my last optimizations

http://codeforces.com/contest/86/submission/15895538

▲ 0 votes • Reply • Permalink



Mike Koltsov 4 Author 5 years ago

Wow, really nice!

▲ 0 votes • Reply • Permalink



shavidze shavidze 5 years ago

mike please, if u can explain me tarjan lca algorithm or give me resources about it ,i search with google tarjan lca algorithm but unfrotunately i can;t understand it , because every explanation is very short and is not clear.

▲ 0 votes • Reply • Permalink



Mike Koltsov / Author 5 years ago

If I was to search for Tarjan explanation, I would definitely visit http://e-maxx.ru/algo/lca\_linear\_offline. It is in russian, though. Maybe Google Translate can help.

▲ 0 votes • Reply • Permalink



shavidze shavidze 5 years ago

thank you very much!

▲ 1 vote • Reply • Permalink



#### Maddela Sai Karthik 5 years ago

Nice

Post mike

Tnxx

▲ 1 vote • Reply • Permalink



#### Vaibhav Tulsyan 5 years ago

Brilliant explanation!

▲ 1 vote • Reply • Permalink



#### Satyarth Agrahari 5 years ago

Thank you so much.. This is an amazing post and it helped me cope up with my fear of questions like this

▲ 1 vote • Reply • Permalink



## Dhruvil Bhuptani 5 years ago

what is the use of blocks here except sorting in "mo order"

▲ 0 votes • Reply • Permalink



#### Mike Koltsov 4 Author 5 years ago

Length of all blocks is chosen in such a way, that we can prove O(N \* sqrt(N)) time complexity. You can vary it and achieve different results.

What else do you expect or need to know about blocks?

▲ 1 vote • Reply • Permalink



#### Dhruvil Bhuptani 5 years ago

@Mike nothing else thank you

▲ 1 vote • Reply • Permalink



### vaibhav goyal @ Edited 4 years ago

Nice work Mike:)

you have done a lot of hard work to prepare this tutorial thank you so much :) and i would like if you will write more on other topics too :)

▲ 1 vote • Reply • Permalink



Saeed Odak 4 years ago

Thanks a lot, really really NICE post. :)

▲ 1 vote • Reply • Permalink



Ankit Kumar 4 years ago

It means that  $k \ge \operatorname{sqrt}(N)$  but the no of elements in a block will always be  $\operatorname{sqrt}(N)$ . Thank you for clarifying my doubt and thank you for such a niece tutorial. It is really helpful for many new programmers. Keep doing good work.

▲ 1 vote • Reply • Permalink



#### Deepankur Gupta 4 years ago

why shouldnt we jump directly to other block when all query of one block is over, why we move in a line to that block,what wil happen if we changes l=r=new\_block\_index; whenerver al query related to previous block is over, As u did in

Our next goal is query [13, 14]. Notice that it starts in different block from the previous query [1, 17]. Consequently, this is the first time mo\_right will move to the left. We need to move mo\_left to 13 and mo\_right to 14.

mo\_right = 16, current\_answer = 112, cnt = [2, 4, 2, 3, 2, 3]

mo\_right = 15, current\_answer = 106, cnt = [2, 4, 1, 3, 2, 3]

mo\_right = 14, current\_answer = 106, cnt = [1, 4, 1, 3, 2, 3]

mo\_left = 2, current\_answer = 99, cnt = [1, 3, 1, 3, 2, 3]

mo\_left = 3, current\_answer = 94, cnt = [1, 2, 1, 3, 2, 3]

mo\_left = 4, current\_answer = 94, cnt = [0, 2, 1, 3, 2, 3]

• 0 votes • Reply • Permalink



#### Mike Koltsov 4 Author 4 years ago

Indeed, it can be done. But then you need to "clear" your data structures explicitly. For example, in our example problem, you need to set `cnt` array to zero. This can be done by simply doing std::fill(cnt, cnt + MAX\_NUMBER, 0) because you do this operation only O(sqrt(N)) times. However, I prefer not to think about how to clear everything - instead, use already written code for moving mo\_left and mo\_right close to each other.

▲ 0 votes • Reply • Permalink



## Deepankur Gupta 4 years ago

okay, and thanks for the lecture. :)

▲ 1 vote • Reply • Permalink



## Shivam Fialok 4 years ago

Brilliant Tutorial on Mo's Algorithm....

Understood in 1 go..!!!

▲ 1 vote • Reply • Permalink



## Priyanshu Varshney 4 years ago

Mike, your notes are amazing. It will be helpful for all if you write one on Heavy-Light decomposition and persistent segment tree. :^)

▲ 1 vote • Reply • Permalink



#### Ashwani Singh 4 years ago

Nice Article

▲ 1 vote • Reply • Permalink



#### madhav gaba 4 years ago

MO's theorum can be applied with updates too. Refer https://www.youtube.com/watch?v=gUpfwVRXhNY&t=1156s

▲ 1 vote • Reply • Permalink



#### suraj prakash narayan 3 years ago

THANKS!!

This tutorial is really helpful. A great explanation.

▲ 1 vote • Reply • Permalink



Vasu Arora 3 years ago

Thank you so much!

▲ 1 vote • Reply • Permalink



#### Manthan Amin 2 years ago

What is the difference between sorting by Left\_Value and sorting by Left\_Value/Block\_Size ? As the Block\_Size is constant!

▲ 0 votes • Reply • Permalink



#### Mike Koltsov 4 Author 2 years ago

Yep, it is constant, yet it makes a lot of sense:)

Dividing by BLOCK\_SIZE helps you to put all Left\_Value-s into "buckets". So, the first BLOCK\_SIZE values go to the 0-th bucket, then the next BLOCK\_SIZE values go to the 1-st bucket, and so on.

And inside the bucket, we sort by Right\_Value.

Inside of one bucket, since we sorted by RIght\_Value, we can move Right\_Value at most to N positions. Which means that in total we won't have more than <number of buckets> \* N moves of Right\_Value.

However, if we skip the BLOCK\_SIZE when sorting, effectively we will have N buckets - i.e. every value is the bucket on its own. That means that Right\_Value can potentially move N  $^*$  N times, which is slow and we want N  $^*$  sqrt(N):)

▲ 0 votes • Reply • Permalink



#### Manthan Amin 2 years ago

I strongly think if we'll sort without without block size logically it will be same, and the algorithm's run time won't change.

proof:

If we'll sort the array with or without dividing by any constant the array will be sorted in regular ascending order.

and we are not using any block specific code in processing queries.

hence the overall effect of logical blocks is nullified.

conclusion, the algorithm will take same time whether or not sorted after dividing by constant.

please reply with counter if any!

▲ 0 votes • Reply • Permalink



#### Mike Koltsov 4 Author 2 years ago

I used to think like that as well, that's why I wrote this article:)

If you think that runtime will be the same, try to get "Accepted" in this task: https://codeforces.com/contest/86/problem/D

As for the counter, please keep in mind that we are sorting \*pairs\*, not just values. So, imagine the case:

(0, 2), (1, 1), (2, 0).

Imagine BLOCK\_SIZE = 3.

Then the ordering with dividing by block size is:

(2, 0), (1,1), (0, 2)

Whereas without dividing it is the same as the original array.

▲ 1 vote • Reply • Permalink



#### Manthan Amin 2 years ago

Ok, got that!

As its floor value division and not division with decimal point my point was

7

wrong.

Thanks:)

▲ 0 votes • Reply • Permalink



## Anuhar Tripathi a year ago

I am still confused how can sorting with block size make all the difference in algorithm's run time. In a question "Dquery" of SPOJ, it is giving TLE when I am sorting without block. Plzz explain.

▲ 0 votes • Reply • Permalink



Mike Koltsov 4 Author 10 months ago

I tried to explain it in the article:(

Try to understand how the time complexity of Mo's algorithm is calculated. Then try to repeat the same method for sorting WITHOUT block. See if it yields the same result (spoiler: no).

▲ 0 votes • Reply • Permalink



Rounik Balida a year ago

nice explained . kudos!!

▲ 1 vote • Reply • Permalink



#### Rajat Bansal 8 months ago

I don't know if you're still out there but this content was worth millions. Loved reading it.

▲ 1 vote • Reply • Permalink



Mike Koltsov 4 Author 8 months ago

Thanks, very nice to hear that:)

▲ 0 votes • Reply • Permalink



Arun Prasad 6 years ago

Very Good Article.

What is the difference between Mo's Algorithm adn Square Root decomposition?

▲ 0 votes • Reply • Permalink



Mike Koltsov 4 Author 6 years ago

My opinion on this matter: Mo's algorithm is one of the tricks of sqrt decomposition. Some people think that Mo's algorithm can not be distinguished from sqrt decomposition, but I think Mo's algorithm is very interesting itself.

▲ 0 votes • Reply • Permalink



#### Abhishek Bind 6 years ago

Very nice Article.

However, I wanted to point out that in the first implementation while processing queries, left and right should be decreased by 1 for correct indexing.

Thanks:)

▲ 0 votes • Reply • Permalink



Mike Koltsov 4 Author 6 years ago

Well, I purposely assume that input queries are 0-based (it is written at the beginning of the example)

▲ 0 votes • Reply • Permalink



Ken Mercado 5 years ago

What is the variable F in the time complexity analysis?

▲ 0 votes • Reply • Permalink

Mike Koltsov 4 Author 5 years ago

?



It is explained in the introduction as

"If we know Func([L, R]), then we can compute Func([L + 1, R]), Func([L - 1, R]), Func([L, R + 1]) and Func([L, R - 1]), each in O(F) time."

F is a function rather than a variable. It is the tightest upper bound you can prove on time complexity of recomputing Func.

▲ 0 votes • Reply • Permalink



#### Priyanshu Varshney 4 years ago

Why is the value of current reduced first, than increased after increasing the value of cnt[x]? Please explain

"// When adding a number, we first nullify it's effect on current

// answer, then update cnt array, then account for it's effect again. "

▲ 0 votes • Reply • Permalink



## Mike Koltsov 4 Author 4 years ago

Suppose cnt[100] = 5. It means that in our current answer there is an addend 5 \* 5 \* 100. We want to turn this addend into 6 \* 6 \* 100. To do that, we subtract 5 \* 5 \* 100, then make 6 out of 5, then add 6 \* 6 \* 100.

▲ 0 votes • Reply • Permalink



## Priyanshu Varshney 4 years ago

Thanks a lot mike

▲ 0 votes • Reply • Permalink



#### Priyanshu Varshney 4 years ago

Also what purpose does "for(const auto &q: queries)" serve

▲ 0 votes • Reply • Permalink



Mike Koltsov 4 Author 4 years ago

That's a C++ "range-based for loop":

http://en.cppreference.com/w/cpp/language/range-for

It is a way of iterating over a container (in this case, vector) of values.

▲ 0 votes • Reply • Permalink



#### Harsh agrawal 3 years ago

can somebody pls tell why my code is printing unexpected output (only sometimes) https://ideone.com/hMdw6l

here is the problem http://www.spoj.com/problems/DQUERY/

▲ 0 votes • Reply • Permalink



#### Amit Kumar 3 years ago

learn this algo from geeksforgeeks , i don't know why it is so elaborated , it is also too much to read , on geeks code is small and easy to understand

▲ 0 votes • Reply • Permalink



Ln (x) dx 3 years ago

Full of unnecessary explanations.

▲ 0 votes • Reply • Permalink



#### Mradul Tiwari a year ago

▲ 0 votes • Reply • Permalink



#### Ganesh Kumar M 9 months ago

Hi All.

Please do checkout my medium article on MO's algorithm https://medium.com/javarevisited/mos-algorithm-range-queries-made-easy-6c35047369ca

▲ 0 votes • Reply • Permalink

### **AUTHOR**



Mike Koltsov ♥ Saint Petersburg, Russia 🖹 3 notes

### TRENDING NOTES

Python Diaries Chapter 3 Map | Filter | Forelse | List Comprehension written by Divyanshu Bansal

Bokeh | Interactive Visualization Library | Use Graph with Django Template written by Prateek Kumar

Bokeh | Interactive Visualization Library | Graph Plotting written by Prateek Kumar

Python Diaries chapter 2 written by Divyanshu Bansal

Python Diaries chapter 1 written by Divyanshu Bansal

more ...

|                         | Resources                              | Solutions                   | Company  | panyService & Support out Us |  |
|-------------------------|--|-----------------------------|----------|------------------------------|--|
|                         | Tech Recruitment                       | Assess Developers           | About Us |                              |  |
|                         | Blog                                   | Conduct Remote              | Press    | Technical Support            |  |
| +1-650-461-4192         | Product Guides                         | Interviews                  | Careers  | Contact Us                   |  |
| contact@hackerearth.com | Developer hiring<br><sup>1</sup> guide | Assess University<br>Talent |          |                              |  |
|                         | -Engineering Blog                      | Organize Hackathons         |          |                              |  |
|                         | Developers Blog                        |                             |          |                              |  |
|                         | Developers Wiki                        |                             |          |                              |  |

Competitive Programming Start a Programming Club

Practice Machine Learning

© 2021 HackerEarth All rights reserved | Terms of Service | Privacy Policy